




PLATFORMS



[Autonomous Machines](#)



[Cloud & Data Center](#)



[Deep Learning & AI](#)

[Design & Pro Visualization](#)

[Healthcare](#)

[High Performance Computing](#)

[Self-Driving Cars](#)

[Gaming & Entertainment](#)

other links

- [Developers](#)
- [Industries](#)
- [Shop](#)
- [Drivers](#)
- [Support](#)
- [About NVIDIA](#)
- [View All Products](#)
- [GPU TECHNOLOGY CONFERENCE](#)
- [NVIDIA Blog](#)
- [Community](#)
- [Careers](#)
- [TECHNOLOGIES](#)

[Newsroom](#)

[NVIDIA in Brief](#)

[Exec Bios](#)

[NVIDIA Blog](#)

[Podcast](#)

[Media Assets](#)

[In the News](#)

[Press Contacts](#)

[Online Press Kits](#)

- [NVIDIA in Brief](#)
- [Exec Bios](#)
- [NVIDIA Blog](#)
- [Podcast](#)
- [Media Assets](#)
- [In the News](#)
- [Press Contacts](#)
- [Online Press Kits](#)

Press Release

Share

[Tweet Twitter](#)

[Share LinkedIn](#)


[Share Facebook](#)

[Email](#)



NVIDIA Vera Rubin Ramps Into Full Production to Power Agentic AI Factories Worldwide

May 31, 2026



News Summary:

- NVIDIA Vera Rubin is ramping into full production, with Taiwan's top server makers and global supply chain leaders manufacturing Vera Rubin-based systems at scale — fueling AI labs, cloud providers and hyperscalers to build tomorrow's intelligence.
- Vera Rubin delivers the POD-scale foundation for next-generation AI factories — with 10x agent throughput at scale compared with the previous-generation NVIDIA Grace Blackwell platform.
- With a proven, open source MGX design, hundreds of NVIDIA supply chain ecosystem partners — 150 in Taiwan alone — across 350+ factories and 30 countries are ramping Vera Rubin.
- Vera Rubin introduces NVIDIA Spectrum-X Ethernet Photonics — now in production — combining co-packaged optics with Spectrum-X switching to enable million-GPU AI factories.

NVIDIA GTC Taipei—NVIDIA today announced the NVIDIA Vera Rubin platform is ramping into full production to power agentic AI factories worldwide.

Taiwan's top server makers and global supply chain leaders are manufacturing Vera Rubin-based systems at scale — fueling AI labs, cloud providers and hyperscalers to build tomorrow's intelligence.

[Vera Rubin](#) delivers NVIDIA's most extensive POD-scale platform — five purpose-built racks operating as one massive AI supercomputer for agentic workloads. The platform unifies NVIDIA Vera Rubin NVL72 systems, NVIDIA Vera CPU, NVIDIA Groq 3 LPX, NVIDIA Vera BlueField-4 STX storage and NVIDIA Spectrum™ -6 SPX Ethernet racks into a fully integrated system. Vera Rubin delivers 10x agent throughput at scale compared with the previous-generation NVIDIA Grace Blackwell platform.

“Agentic AI is a new kind of workload. One prompt can launch a thousand-step journey of reasoning, retrieval, tool use and response generation,” said Jensen Huang, founder and CEO of NVIDIA. “Vera Rubin was built for this moment — an AI factory engine that delivers intelligence at scale, with the performance, efficiency and security needed to power the next industrial revolution.”

Vera Rubin Ramp

Vera Rubin marks the third generation of NVIDIA MGX™ rack-scale systems. With a proven, open source MGX design, hundreds of NVIDIA supply chain ecosystem partners — 150 in Taiwan alone — across 350+ factories and 30 countries are ramping Vera Rubin.

Top system builders, infrastructure software and storage partners are in full-scale production of Vera Rubin. This includes Dell Technologies, HPE, Lenovo and Supermicro, as well as AIC, Aivres, ASRock

Rack, [ASUS](#), Cloudian, [Compal](#), DDN, Everpure, Foxconn, GIGABYTE, Hitachi Vantara, Hyve Solutions, IBM, Inventec, MinIO, MiTAC Computing, [MSI](#), NetApp, Nutanix, Pegatron, Quanta Cloud Technology (QCT), VAST Data, WEKA, Wistron and Wiwynn.

Building the Fabric for Million-GPU AI Factories

To support scale-out and scale-across AI factory deployments, the Vera Rubin platform introduces NVIDIA Spectrum-X™ Ethernet Photonics, the world's first co-packaged-optics (CPO)-based switches with 200Gb/s SerDes — now in production.

Spectrum-X Ethernet Photonics, a new generation of switching technology built on CPO, delivers 5x better power efficiency, 5x longer AI uptime and 1.3x faster time to deployment than networks using traditional transceivers.

By simplifying design and freeing more power for compute, NVIDIA co-packaged optics networking provides the foundational fabric for million-GPU AI factories, with CoreWeave, [Lambda](#) and Oracle Cloud Infrastructure among the first ecosystem partners and adopters.

The NVIDIA Vera Rubin platform also integrates NVIDIA BlueField®-4 DPUs, featuring software-defined networking at speeds of up to 800Gb/s and built-in multi-tenant isolation. With the NVIDIA BlueField-4 Advanced Secure Trusted Resource Architecture, customers can simplify network operations, improve tenant isolation and gain greater control across million GPU AI clusters.

Secure AI for AI Factories

AI factories are increasingly processing proprietary data, regulated content and mission-critical models in agentic workflows. This requires advanced infrastructure security tailored to autonomous agents in shared or cloud environments where infrastructure cannot be implicitly trusted.

The Vera Rubin platform was designed with full-stack NVIDIA Confidential Computing for a trusted execution environment at rack scale. Vera Rubin NVL72 combines Vera CPUs, Rubin GPUs, NVIDIA NVLink™ networking and security features into a unified platform, encrypting data across high-speed interconnects. This provides hardware-level attestation to ensure the system is tamper-proof.

Cloud providers CoreWeave, Firmus, GMI Cloud, IBM Cloud, IREN, Lambda, Microsoft Azure, Nebius, Nscale, SpaceXAI and Vultr are adopting NVIDIA Confidential Computing.

Delivering this level of protection at POD scale also requires a programmable software layer capable of enforcing, orchestrating and adapting security policies across the entire system. The NVIDIA DOCA™ software platform delivers advanced security across every Vera Rubin platform rack and layer of the AI factory — protecting data, agents, context memory and AI inference through capabilities enforced directly in BlueField-4 silicon.

DOCA enables multi-tenant network isolation, zero-trust policy enforcement, runtime threat detection and end-to-end encryption at speeds of up to 800Gb/s, all without taxing host CPU resources, so enterprises can scale AI factories with confidence.

Accelerating the Buildout of AI Factories

The NVIDIA DSX™ platform provides the complete design and operational foundation for Vera Rubin AI factories — unifying reference designs, simulation, infrastructure software, facilities and ecosystem technologies to help build and operate energy-efficient AI factories optimized for lowest token cost.

Built for the Vera Rubin POD architecture, DSX aligns every layer of the stack — from silicon and systems to lifecycle management and multi-tenant operations — dramatically accelerating deployment and setting a new bar for operational reliability and resiliency at scale.

Dell Technologies, HPE, Lenovo and Supermicro together with ASUS, Foxconn, GIGABYTE, Pegatron, Quanta Cloud Technology (QCT), Wistron and Wiwynn are adopting NVIDIA DSX to accelerate AI factory ramp with Vera Rubin.

Availability

Production shipments of Vera Rubin are set to begin starting this fall.

Watch Huang's [keynote](#) and learn more at [NVIDIA GTC Taipei](#).

About NVIDIA

[NVIDIA](#) (NASDAQ: NVDA) is the world leader in AI and accelerated computing.

Certain statements in this press release including, but not limited to, statements as to: expectations with respect to growth, performance, availability, and benefits of NVIDIA's products, services and technologies, and related trends and drivers; expectations with respect to NVIDIA's third party arrangements, including with its collaborators and partners; expectations with respect to technology developments, and related trends and drivers; projected market growth and trends; expectations with respect to AI and related industries; and other statements that are not historical facts are forward-looking statements within the meaning of Section 27A of the Securities Act of 1933, as amended, and Section 21E of the Securities Exchange Act of 1934, as amended, which are subject to the "safe harbor" created by those sections based on management's beliefs and assumptions and on information currently available to management and are subject to risks and uncertainties that could cause results to be materially different than expectations. Important factors that could cause actual results to differ materially include: global economic and political conditions; NVIDIA's reliance on third parties to manufacture, assemble, package and test NVIDIA's products; the impact of technological development and competition; development of new products and technologies or enhancements to NVIDIA's existing products and technologies; market acceptance of NVIDIA's products or NVIDIA's partners' products; design, manufacturing or software defects; changes in consumer preferences or demands; changes in industry standards and interfaces; unexpected loss of performance of NVIDIA's products or technologies when integrated into systems; NVIDIA's ability to realize the potential benefits of business investments or acquisitions; and changes in applicable laws and regulations, as well as other factors detailed from time to time in the most recent reports NVIDIA files with the Securities and Exchange Commission, or SEC, including, but not limited to, its Annual Report on Form 10-K and Quarterly Reports on Form 10-Q. Copies of reports filed with the SEC are posted on the company's website and are available from NVIDIA without charge. These forward-looking statements are not guarantees of future performance and speak only as of the date hereof, and, except as required by law, NVIDIA disclaims any obligation to update these forward-looking statements to reflect future events or circumstances.

Many of the products and features described herein remain in various stages and will be offered on a when-and-if-available basis. The statements above are not intended to be, and should not be interpreted as a commitment, promise, or legal obligation, and the development, release, and timing of any features or functionalities described for our products is subject to change and remains at the sole discretion of NVIDIA. NVIDIA will have no liability for failure to deliver or delay in the delivery of any of the products, features or functions set forth herein.

© 2026 NVIDIA Corporation. All rights reserved. NVIDIA, the NVIDIA logo, BlueField, NVIDIA DOCA, NVIDIA DSX, NVIDIA MGX, NVIDIA Spectrum and NVIDIA Spectrum-X are trademarks and/or registered trademarks of NVIDIA Corporation in the U.S. and/or other countries. Other company and product names may be trademarks of the respective companies with which they are associated.


Media Contacts

Kristin Uchiyama
Corporate Communications
NVIDIA Corporation
press@nvidia.com

Downloads


- [Download Press Release](#)
- [Download Attachments](#)

More Images



•





More News




[NVIDIA Sets Conference Call for Second-Quarter Financial Results](#)

July 29, 2026



[Ilya Sutskever's Safe Superintelligence Inc. and NVIDIA Announce Long-Term Strategic Partnership](#)

July 27, 2026



[NVIDIA Expands NVIDIA Agent Toolkit With NVIDIA PhysicsNeMo and CUDA-X Libraries to Transform How the World Engineers, Designs and Builds](#)

July 26, 2026



[SK Group and NVIDIA Expand Strategic Partnership Across AI Factories and Next-Generation Memory](#)

July 24, 2026



[NAVER, NVIDIA and Brookfield to Expand Korea's National AI Factory Infrastructure Buildout](#)

July 24, 2026

About NVIDIA

[NVIDIA](#) (NASDAQ: NVDA) is the world leader in AI and accelerated computing.

Certain statements in this press release including, but not limited to, statements as to: expectations with respect to growth, performance, availability, and benefits of NVIDIA's products, services and technologies, and related trends and drivers; expectations with respect to NVIDIA's third party arrangements, including with its collaborators and partners; expectations with respect to technology developments, and related trends and drivers; projected market growth and trends; expectations with respect to AI and related industries; and other statements that are not historical facts are forward-looking statements within the meaning of Section 27A of the Securities Act of 1933, as amended, and Section 21E of the Securities Exchange Act of 1934, as amended, which are subject to the "safe harbor" created by those sections based on management's beliefs and assumptions and on information currently available to management and are subject to risks and uncertainties that could cause results to be materially different than expectations. Important factors that could cause actual results to differ materially include: global economic and political conditions; NVIDIA's reliance on third parties to manufacture, assemble, package and test NVIDIA's products; the impact of technological development and competition; development of new products and technologies or enhancements to NVIDIA's existing products and technologies; market acceptance of NVIDIA's products or NVIDIA's partners' products; design, manufacturing or software defects; changes in consumer preferences or demands; changes in industry standards and interfaces; unexpected loss of performance of NVIDIA's products or technologies when integrated into systems; NVIDIA's ability to realize the potential benefits of business investments or acquisitions; and changes in applicable laws and regulations, as well as other factors detailed from time to time in the most recent reports NVIDIA files with the Securities and Exchange Commission, or SEC, including, but not limited to, its Annual Report on Form 10-K and Quarterly Reports on Form 10-Q. Copies of reports filed with the SEC are posted on the company's website and are available from NVIDIA without charge. These forward-looking statements are not guarantees of future performance and speak only as of the date hereof, and, except as required by law, NVIDIA disclaims any obligation to update these forward-looking statements to reflect future events or circumstances.

Many of the products and features described herein remain in various stages and will be offered on a when-and-if-available basis. The statements above are not intended to be, and should not be interpreted as a commitment, promise, or legal obligation, and the development, release, and timing of any features or functionalities described for our products is subject to change and remains at the sole discretion of NVIDIA. NVIDIA will have no liability for failure to deliver or delay in the delivery of any of the products, features or functions set forth herein.

© 2026 NVIDIA Corporation. All rights reserved. NVIDIA, the NVIDIA logo, BlueField, NVIDIA DOCA, NVIDIA DSX, NVIDIA MGX, NVIDIA Spectrum and NVIDIA Spectrum-X are trademarks and/or registered trademarks of NVIDIA Corporation in the U.S. and/or other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

Media Contacts

Global contacts for media inquiries.

[All Contacts](#)

Stay Informed

Newsroom updates delivered to your inbox.

[Subscribe](#)

Corporate Information

- [About NVIDIA](#)
- [Corporate Overview](#)
- [Technologies](#)
- [NVIDIA Research](#)
- [Investors](#)
- [Social Responsibility](#)
- [NVIDIA Foundation](#)

Get Involved

- [Forums](#)
- [Careers](#)
- [Developer Home](#)
- [Join the Developer Program](#)
- [NVIDIA Partner Network](#)
- [NVIDIA Inception](#)
- [Resources for Venture Capitalists](#)
- [Venture Capital \(NVentures\)](#)
- [Technical Training](#)
- [Training for IT Professionals](#)
- [Professional Services for Data Science](#)

News & Events

- [Newsroom](#)
- [NVIDIA Blog](#)
- [NVIDIA Technical Blog](#)
- [Webinars](#)
- [Stay Informed](#)
- [Events Calendar](#)


- [NVIDIA GTC](#)
- [NVIDIA On-Demand](#)




Sign Up for NVIDIA News

[Subscribe](#)

Follow NVIDIA





nVIDIA®

[USA - United States](#)

- [Privacy Policy](#)
- [Your Privacy Choices](#)
- [Manage Settings](#)
- [Terms of Service](#)
- [Accessibility](#)
- [Corporate Policies](#)
- [Product Security](#)
- [Contact](#)

Copyright © 2026 NVIDIA Corporation